

## Write-Up

Anjin- Worked on pipeline functions and LRU functions, also debugging and code cleanup

Dhruv- Worked on the LRU hit and miss functions (not implemented versions), performance assessment and write up

Hao- Worked on branch prediction and implementing the `push_pipeline_stage` function

Jay- Worked on cache data structure and associated methods, some editing

In our project we implemented several methods which allowed our pipeline to run accurately. First we ask the user to enter in a cache index, block size and associativity. We then make sure that we don't not exceed the cache size of 10240. Next we dynamically create the cache and initialize the pipeline. We then go through and parse the information in our trace file.

Every time we parse one line, we check to see if the address is already in our cache, using `iplc_sim_trap_address`. If it is in our cache, then we call the `iplc_sim_LRU_update_on_hit` method which will simply add to our counter for the amount of cache hits and update the value in our cache. However, if the address is not in our cache then we call the `iplc_sim_LRU_replace_on_miss` which goes through and looks to find either an empty spot in the cache or looks to find the lowest value in the cache and replace the old data with the new address's data. We then also update the amount of cache misses. If we do have a miss, in our `iplc_sim_parse_instruction` method we push the instruction through the pipeline. The purpose of doing this is that we do not double count cycles but instead have them overlap.

In the pipeline simulator, `iplc_sim_push_pipeline_stage`, we go through seven stages. The first stage is the writeback stage in we write our results into the register file. The second stage we check to see if the branch prediction was right. The third stage is to check for any LW delays. If there are any then we add 9 cycles to our cycle count. We are adding one additional cycle later one so that is why we do not add the full ten cycle delay. In the fourth stage we look for a SW memory access and data miss and add 9 cycles if it does happen. In the fifth date we just implement the total cycles by one. In the sixth stage we push through onto the next stage of the pipeline. In the seventh stage we reset the fetch stage to NOP.


We then go back into the `iplc_sim_parse_instruction` and parse the instruction. We check to see if then instruction if "add", "sll" or "ori" as if it is then we assign values to our `dest_reg`, `src_reg` and `src_reg`. We then call the `iplc_sim_process_pipeline_rtype` method which basically assigns the `dest_reg`, `src_reg` and `src_reg` to the Fetch in the pipeline. Next, we check to see if the instruction is "lui", if it is we then call the `iplc_sim_process_pipeline_rtype` method again but this time we would assign the same value to `dest_reg` as the last conditional but assign the values -1 to both `src_reg` and `src_reg2`. We then check to see if the instruction is "lw" or "sw", if it is "lw" we then set the `dest_reg` value and call the `iplc_sim_process_pipeline_lw` method which essentially sets the values of the `dest_reg`, `base_reg` and `data_address` for the lw in the pipeline. If the instruction was "sw" however, then we assign the proper value to the `src_reg` and call,

iplc\_sim\_process\_pipeline\_sw. If the instruction is “beq” then we call iplc\_sim\_process\_pipeline\_branch. If the instruction was “jal”, “jr” or “j” then we call iplc\_sim\_process\_pipeline\_jump. If the instruction is “syscall” we call iplc\_sim\_process\_pipeline\_syscall. If the instruction is “nop”, then we call iplc\_sim\_process\_pipeline\_nop. All of these methods just assign values to the targeted pipeline. Lastly, we call iplc\_sim\_finalize which will ensure all instructions are done proceeded and print important statistics.

### Performance Assessment


Based on the results of our 18 different tests. We have concluded that the best highest performing configuration is the one with a cache index, Block Size and Associativity of 4 with the branches not taken. This can be proved by 3 key statistics.

First, we will talk about CPI. CPI is the average number of CPU clock cycles that occur per an instruction are being executed. It can be calculated by taking the total amount of cycles and dividing it by the amount of instructions.




As we can see the 9<sup>th</sup> bar on the bar graph is below the average. The average CPI for all 18 tests was 1.248122389 while the CPI of this specific configuration is 1.057491. The standard deviation for this data was 0.135008803. Thus, we know that the CPI of 9<sup>th</sup> bar is well below one standard deviation meaning that's its CPI is in about the lowest 13 percentile. A low CPI is good because you want the least amount of cycles to pass per instruction, meaning that it takes less time to complete and instruction.

Second, we will discuss the miss rate. The miss rate is fraction of memory references not found in the cache. It can be calculated by dividing the cache misses by the cache accesses.



We can see that the 9<sup>th</sup> configuration had an extremely low miss rate to many of the other configurations. The average miss rate was 0.015827444 while the miss rate for the 9<sup>th</sup> configuration was 0.002175. The standard deviation of the data is 0.012806946, thus the 9<sup>th</sup> configuration is more than a standard deviation away from the average meaning that again, it is in the lowest 13 percentile. Having a low miss rate is good because that means that you spend less time copying things from the main memory to the cache.

Lastly, we would like to bring up the statistic of the total correct branch predictions. This basically means how many times did our program guess the right way to branch off to.



This data clearly shows that the branch not taken method usually yielded better result than the branch taken method by a lot.

These three statistic give us plentiful evidence that the 9<sup>th</sup> configuration performs because it has one of the lowest CPI's and miss rates and has a higher correct branch prediction rate than the branch taken method configurations.

Here is the full data chart:

Configur ation Number	Cache Index	Block Size	Levels of Associati vity	Total Cycles	Total Instructi on	Total Branch Instructi ons	Total Correct Branch Predictio ns	CPI	Number of Cache Accesses	Number of Cache Misses	Number of Cache Hits
1	7	1	1	48547	34753	7044	5730	1.396915	35863	1390	3447
2	6	1	2	40556	34753	7044	5739	1.166978	35863	502	3536
3	5	1	4	39309	34753	7044	5744	1.131097	35863	363	3550
4	6	2	1	47757	34753	7044	5276	1.374184	35863	1301	3456
5	5	2	2	39261	34753	7044	5736	1.129715	35863	357	3550
6	4	2	4	37848	34753	7044	5746	1.089057	35863	200	3566
7	6	4	1	42979	34753	7044	5729	1.236699	35863	770	3509
8	5	4	2	37345	34753	7044	5749	1.074583	35863	144	3571
9	4	4	4	36751	34753	7044	5749	1.057491	35863	78	3578
10	7	1	1	52998	34753	7044	1279	1.524991	35863	1390	3447
11	6	1	2	45015	34753	7044	1280	1.295284	35863	502	3536
12	5	1	4	43769	34753	7044	1284	1.259431	35863	363	3550
13	6	2	1	52193	34753	7044	1290	1.501827	35863	1301	3456
14	5	2	2	43707	34753	7044	1290	1.257647	35863	357	3550
15	4	2	4	42304	34753	7044	1290	1.217276	35863	200	3566
16	6	4	1	47417	34753	7044	1291	1.3644	35863	770	3509
17	5	4	2	41803	34753	7044	1291	1.20286	35863	144	3571
18	4	4	4	41209	34753	7044	1291	1.185768	35863	78	3578